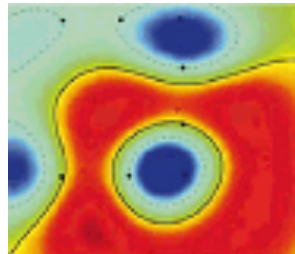
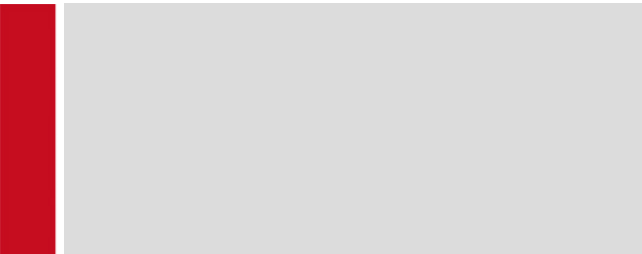




WiSe 2023/24

Machine Learning 1/1-X



Lecture 9

Neural Networks 2

Outline

- ▶ Recap:
 - ▶ The Perceptron
 - ▶ Artificial Neural Networks
 - ▶ Forward / Backward Propagation
- ▶ Optimizing Neural Networks:
 - ▶ Bad Local Minima and Pathological Curvature
 - ▶ Initialization / Centering / Momentum
- ▶ Regularizing Neural Networks:
 - ▶ Hinge Loss
 - ▶ Perturbations

Recap: The Perceptron



F. Rosenblatt (1928–1971)

- ▶ Proposed by F. Rosenblatt in 1958.
- ▶ Classifier that perfectly separates training data (if the data is linearly separable).
- ▶ Trained using a simple and cheap iterative procedure.
- ▶ The perceptron gave rise to artificial neural networks.

Recap: The Perceptron Algorithm

Consider the linear model $y = \mathbf{w}^\top \mathbf{x} + b$, where $\text{sign}(y)$ gives the classification decision. Let $t \in \{-1, +1\}$ be the binary class label associated to \mathbf{x} .

Algorithm


- ▶ Iterate over all data points $\{(\mathbf{x}_k, t_k); k = 1 \dots, N\}$ (multiple times).
 - ▶ Compute $y_k = \mathbf{w}^\top \mathbf{x}_k + b$
 - ▶ If \mathbf{x}_k is correctly classified (i.e. $\text{sign}(y_k) = t_k$), continue.
 - ▶ If \mathbf{x}_k is wrongly classified (i.e. $\text{sign}(y_k) \neq t_k$), apply:

$$\mathbf{w} \leftarrow \mathbf{w} + \gamma \cdot \mathbf{x}_k t_k \quad \text{and} \quad b \leftarrow b + \gamma \cdot t_k$$

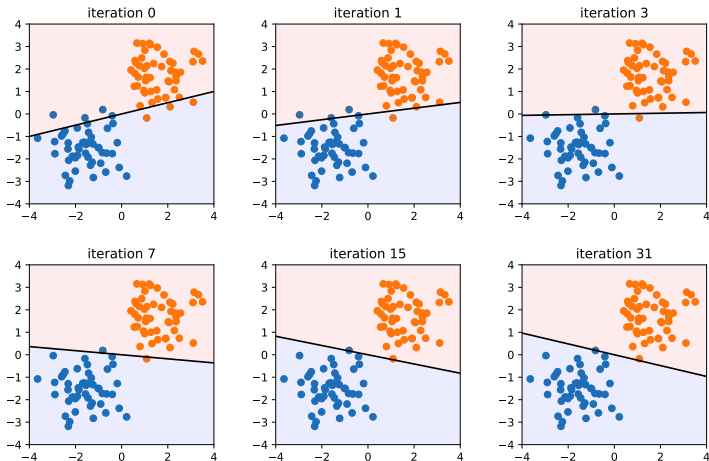
- ▶ Stop once all examples are correctly classified.

The perceptron can also be seen as a stochastic gradient descent of the error function

$$\mathcal{E}(\mathbf{w}, b) = \frac{1}{N} \sum_{k=1}^N \underbrace{\max(0, -y_k t_k)}_{\mathcal{E}_k(\mathbf{w}, b)}$$



Recap: Perceptron at Work



Recap: Nonlinear Classification

Observation:


- ▶ The perceptron builds a decision function which is linear in input space.
- ▶ In practice, the data may not be linearly separable.

Key Idea:

- ▶ Transform the data nonlinearly through some function ϕ before applying the linear decision function.


$$f(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

- ▶ *Example:* $\phi(\mathbf{x}) = [x_1, x_2, x_1^2, x_2^2, x_1x_2]$ and $\mathbf{w} \in \mathbb{R}^5$.




Recap: Artificial Neural Networks

- ▶ Models that are inspired by the way the brain represents sensory input and learn from repeated stimuli.
- ▶ Neuron activations can be seen as a nonlinear transformation of the sensory input (similar to $\phi(x)$).
- ▶ The neural representation adapts itself after repeated exposure to certain stimuli (plasticity).




Recap: The Artificial Neuron



- ▶ Simple multivariate, nonlinear and differentiable function.
- ▶ Ultra-simplification of the biological neuron that retains two key properties: (1) ability to produce complex nonlinear representations when many neurons are interconnected (2) ability to learn from the data.

Recap: Artificial Neural Networks



$$z_3 = x_1 w_{13} + x_2 w_{23} \quad a_3 = g(z_3)$$

$$z_4 = x_2 w_{24} + a_3 w_{34} \quad a_4 = g(z_4)$$


$$y = a_3 \cdot v_3 + a_4 v_4$$

Recap: Training a Neural Network

Idea:

- ▶ Use the same error function as the perceptron, but replace the perceptron output by the neural network output:

$$\mathcal{E}(\theta) = \frac{1}{N} \sum_{k=1}^N \underbrace{\max(0, -y_k t_k)}_{\mathcal{E}_k(\theta)}$$



and compute the gradient of the error function w.r.t. the parameters θ of the neural network.

Question:

- ▶ How to compute the gradient of the error function?

Recap: Error Backpropagation

Idea:

- ▶ The gradient can be expressed in terms of gradient in the higher layers using the multivariate chain rule.

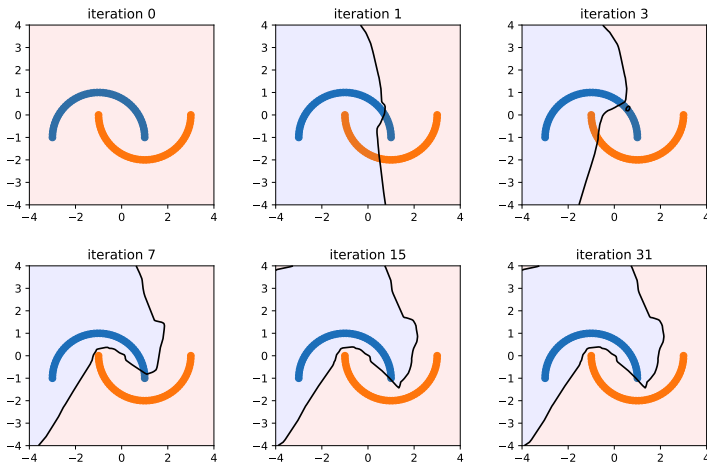
$$\frac{\partial \mathcal{E}}{\partial z_i} = \sum_j \frac{\partial z_j}{\partial z_i} \frac{\partial \mathcal{E}}{\partial z_j}$$

- ▶ Similarly, one can then extract the gradient w.r.t. the parameters of the model as:

$$\frac{\partial \mathcal{E}}{\partial w_{ij}} = \frac{\partial z_j}{\partial w_{ij}} \frac{\partial \mathcal{E}}{\partial z_j}$$

- ▶ Gradients can be computed one after the other in a message passing fashion in **O(forward pass)**. The algorithm is known as *error backpropagation*.

Recap: Neural Network at Work



Today's Lecture

Part 1:

- ▶ Optimizing Neural Networks
(making sure neural networks can be trained effectively and efficiently).


Part 2:

- ▶ Regularizing Neural Networks
(making sure the learned models are good and generalize well).

Part 1: Optimizing Neural Networks

Questions:


- ▶ How hard is it to optimize a neural network?
- ▶ Are we guaranteed to converge to a good local minimum?
- ▶ How quickly do we converge to this local minimum?



Neural Networks: Initialization

Experiment:

- ▶ Train the network with different weight initializations:



Observations:

- ▶ Some runs do not converge to a solution that perfectly classifies the data (e.g. when weights are initialized to zero, training is stuck on a plateau corresponding to a constant classifier).

某些训练运行无法收敛到一个能够完美分类数据的解决方案（例如，当权重初始化为零时，训练会卡在对应于常数分类器的平坦区域上）。

Neural Networks: Initialization

权重应随机初始化（这样可以打破参数空间中的对称性，并减少落入糟糕的局部极小值或陷入平坦区域的风险）。

如果有必要，可以用不同的随机种子多次训练网络，并保留误差最小的网络。

- ▶ Weights should be initialized *randomly* (it breaks symmetries in parameter space and reduces the risk of landing in bad local minima or getting stuck on plateaus).
- ▶ If necessary, train the network *multiple times* using different random seeds, and retain the network with the *lowest error*.

He-et-al¹ random initialization heuristic:

For neural networks with ReLU neurons, it is recommended to use:

$w_{ij} \sim \mathcal{N}(0, \sigma^2)$ with $\sigma = \sqrt{2/\# \text{ input connections}}$ and where \mathcal{N} is the Gaussian distribution.

Example: For a two-layer neural network with 50 input features, 200 hidden neurons, and 1 output, weights in the first layer should be initialized using $\sigma = \sqrt{2/50}$, and in the second layer using $\sigma = \sqrt{2/200}$.

¹He et al. (2015) Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification


Neural Networks: Pathological Curvature

神经网络：病态曲率

另一个问题：

由于误差函数的病态曲率，神经网络可能会收敛到某个较好的局部最优解，但收敛速度非常缓慢。

Another problem: The neural network may converge to some good local optimum but slowly due to *pathological curvature* of the error function.



Source: Martens. Deep Learning via Hessian-free Optimization U Toronto, 2010.

Neural Networks: Pathological Curvature



可以通过观察误差函数的海森矩阵 (Hessian Matrix) 的特征值来描述误差函数的曲率:

The curvature of the error function can be characterized by looking at the eigenvalues of the Hessian of the error function:



$$H = \left(\frac{\partial^2 \mathcal{E}(\theta)}{\partial \theta_i \partial \theta_j} \right)_{ij} \quad \lambda_1, \dots, \lambda_{|\theta|} = \text{eigvals}(H)$$

The higher the ratio between the highest and the lowest eigenvalue (aka. condition number), the slower the convergence.

Two-dimensional example:



Neural Networks: Pathological Curvature




The lower the condition number, the better.

λ_1 is the largest eigenvalue

Hessian Analysis (Linear Case, MSE)

Consider the simplest case of a homogeneous linear model $y = \mathbf{w}^\top \mathbf{x}$. Consider $\mathcal{E}(\mathbf{w}) = \mathbb{E}[0.5 \cdot (1 - y \cdot t)^2]$, the **mean square error (MSE)**, which is easier to analyze than the **perceptron loss**.



The Hessian of the error function is given by:

$$H = \frac{\partial^2 \mathcal{E}}{\partial \mathbf{w}^2} = \mathbb{E}[\mathbf{x} \mathbf{x}^\top]$$

i.e. the data uncentered covariance.

Condition Number: Assuming $\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \sigma^2 I)$, the Hessian reduces to $H = \sigma^2 I + \boldsymbol{\mu} \boldsymbol{\mu}^\top$, and the condition number is

$$\frac{\lambda_1}{\lambda_d} = 1 + \frac{\|\boldsymbol{\mu}\|^2}{\sigma^2}.$$

(show this in the homework). In other words, the condition number can be reduced by **centering the data** before training.

Hessian Analysis (General Case, MSE)

- Hessian of a neural network (LeCun et al. 1998): $\frac{\partial \mathcal{E}}{\partial \mathbf{Y}} \left(\frac{\partial \mathbf{Y}^\top}{\partial \boldsymbol{\theta}} \frac{\partial \mathcal{E}}{\partial \mathbf{Y}} \frac{\partial \mathbf{Y}}{\partial \boldsymbol{\theta}} + \frac{\partial^2 \mathcal{E}}{\partial \boldsymbol{\theta}^2} \right)$

$$H = \frac{\partial^2 \mathcal{E}}{\partial \boldsymbol{\theta}^2} = \frac{\partial \mathbf{Y}^\top}{\partial \boldsymbol{\theta}} \frac{\partial^2 \mathcal{E}}{\partial \mathbf{Y}^2} \frac{\partial \mathbf{Y}}{\partial \boldsymbol{\theta}} + \frac{\partial \mathcal{E}}{\partial \mathbf{Y}} \frac{\partial^2 \mathbf{Y}}{\partial \boldsymbol{\theta}^2} \quad (1)$$

where \mathbf{Y} is the vector containing the predictions for all data points.

- Consider the part of the Hessian relating parameters of a specific neuron k . From Eq. (1), and assuming MSE, we get:

$$[H_k]_{jj'} = \frac{\partial^2 \mathcal{E}}{\partial w_{jk} \partial w_{j'k}} = \mathbb{E}[a_j a_{j'} \delta_k^2] + \mathbb{E}[a_j \frac{\partial \delta_k}{\partial w_{j'k}} \cdot (y - t)]$$


where δ_k is a shortcut notation for $\partial \mathcal{E} / \partial z_k$.

- Assuming $a_j a_{j'}$ and δ_k^2 to be independent, and δ_k insensitive to $w_{j'k}$, we get:

$$H_k \approx \mathbb{E}[\mathbf{a} \mathbf{a}^\top] \cdot \mathbb{E}[\delta_k^2]$$

which has a similar structure as in the analysis for the linear case.

- This motivates **centering activations** as well.



Hessian Analysis: Takeaway Message

为了降低条件数，并从而简化神经网络的优化过程：

使输入数据居中，并确保每一层的激活值也居中。

注意：


在实际中精确地使激活值居中可能很困难，但可以通过选择满足 $g(0) = 0$ 的激活函数来实现大致的居中。

To lower the condition number and consequently ease optimization of the neural network:

Center the input data, and also activations at each layer.

Note: Exactly centering activations can be difficult in practice, however, reasonably centered activations can be promoted by choosing an activation function satisfying $g(0) = 0$.

- ▶ Examples of **suitable** activation functions (where $g(0) = 0$) are the *ReLU* ($\max(0, z)$), the *centered softplus* ($\log(1 + \exp(z)) - \log(2)$), or the *hyperbolic tangent* ($\tanh(z)$).
- ▶ Examples of **less suitable** activation functions (where $g(0) \neq 0$) are the *standard softplus* ($\log(1 + \exp(z))$), or the *logistic sigmoid* ($\exp(z)/(1 + \exp(z))$).




Improving Gradient Descent

改进梯度下降

- 居中可能显著减少误差函数的病态曲率，但不能完全消除它（尤其是对于深度神经网络而言）。
- **一种补充方法：可以通过修改优化方法，使其沿低曲率方向加速移动。
- 这可以通过在梯度下降中引入**动量（momentum）**来实现。

- ▶ Centering may significantly reduce the pathological curvature, but does not completely eliminate it (especially for deep neural networks).
- ▶ A complementary approach to centering is to modify the optimization procedure to move faster along directions of low curvature.
- ▶ This can be achieved by introducing a '*momentum*' in the gradient descent:



Improving Gradient Descent

(Stochastic) Gradient Descent:

$$\theta = \theta - \gamma \cdot \frac{\partial \mathcal{E}_k}{\partial \theta} \quad \text{with } k \sim \{1, \dots, N\}$$

where γ is the learning rate.

(Stochastic) Gradient Descent + Momentum:

$$\Delta_{\text{new}} = \mu \Delta_{\text{old}} + \frac{\partial \mathcal{E}_k}{\partial \theta} \quad \text{with } k \sim \{1, \dots, N\}$$
$$\theta = \theta - \gamma \Delta_{\text{new}}$$

where γ is the learning rate and μ is the momentum.


Heuristics:

- ▶ Typical learning rate and momentum: $\gamma = 0.01$ and $\mu = 0.9$ (need to be reduced if training diverges).

Part 2: Regularizing Neural Networks

Questions:

- ▶ Do neural networks trained with gradient descent and the perceptron loss yield good models?
- ▶ How to help neural networks to learn solutions that generalize well to new data?




Learning Well-Generalizing Solutions

So far:

- ▶ We have used the perceptron loss:

$$\mathcal{E}_k(\theta) = \max(0, -y_k t_k)$$

which becomes zero as soon as the current data point is being correctly classified.




Idea: 对离决策边界过近的点（即添加一个间隔）施加惩罚，也就是 Hinge 损失（Hinge Loss）：


- ▶ Impose a penalty to points that lie too close to the decision boundary (i.e. add a margin), aka. the **Hinge Loss**:

$$\mathcal{E}_k(\theta) = \max(0, 1 - y_k t_k)$$


This is similar to the slack variables penalties in the soft-margin SVM formulation.



这类似于软间隔支持向量机（SVM）公式中的松弛变量惩罚项。



Hinge Loss at Work



Observation:

- ▶ Using the hinge loss, the decision function of the neural network moves away from the data.
- ▶ This leads to a better generalization performance.

More Loss Functions


Name	Formula	Margin	Outliers ²
Perceptron loss	$\max(0, -y \cdot t)$	no	yes
Hinge loss	$\max(0, 1 - y \cdot t)$	yes	yes
Squared hinge loss	$\max(0, 1 - y \cdot t)^2$	yes	no
Log-loss	$\log(1 + \exp(-y \cdot t))$	yes	yes

²Indicates whether the loss function is robust to outliers (e.g. mislabelings).

What is Still Missing?

Problem:

- ▶ If multiplying the weights by some factor, the hinge loss can be trivially reduced until it reaches zero, without actually changing the decision boundary.



- ▶ In the SVM, this was avoided by constraining $\|\mathbf{w}\|^2$ to be small.

Idea:


- ▶ To arrive at a well-generalizing model, we also need to regularize the neural network itself.

Regularizing the Neural N


- **权重衰减 (Weight Decay) :**
在目标函数中添加一项 $\lambda \cdot \|\theta\|^2$ 。
这相当于在每次训练迭代中通过某个常数因子来缩小权重。
- **数据扰动 (Data Perturbation) :**
对输入样本 (例如, 添加高斯噪声) 和标签 (随机翻转) 应用随机扰动。
- **表示扰动 (Representation Perturbation) :**
Dropout 方法 (Srivastava, 2014) 将数据扰动的概念扩展到中间表示层, 训练时随机关闭和开启神经元。
- **训练噪声 (Training Noise) :**
已发现随机梯度下降 (SGD) 与标准梯度下降相比具有正则化效果。

Many approaches have been proposed and are often combined in practice. Some of the most popular ones are:

- ▶ **Weight decay:** Add a term $\lambda \cdot \|\theta\|^2$ to the objective. This is equivalent to multiplying the weights at each training iteration by some constant factor.
- ▶ **Data perturbation:** Apply random perturbations to the input examples (e.g. Gaussian additive noise), and to the labels (random flips).
- ▶ **Representation perturbation:** The Dropout method (Srivastava'14) extends the idea of data perturbation to intermediate representations, by randomly turning on and off neurons while training.
- ▶ **Training noise:** Stochastic gradient descent was found to have a regularizing effect compared to standard gradient descent.



Data Perturbation at Work



Observation:

- ▶ Data perturbations pull the decision boundary further away from the class manifolds, and typically further improves generalization.

数据扰动将决策边界进一步拉离类别流形（class manifolds），并通常进一步提高模型的泛化能力。



Neural Networks vs. SVMs/Kernels

SVMs/Kernels:

- ▶ Relatively straightforward to optimize and regularize.
- ▶ Doesn't scale well to very large datasets with millions of examples.

Neural Networks:

- ▶ Can be trained on very large distributions (with GPUs).
- ▶ Challenging to optimize and regularize neural networks. Lots of tricks and heuristics involved. Can make results harder to reproduce and less transparent.




, etc.

Neural Networks: Further Topics

Further Topics:

- ▶ **Incorporate prior knowledge:** Data extension, weight sharing, unsupervised pretraining, self-supervised pretraining, transfer learning, multitask learning.



- ▶ **Predicting structured data:** Convolutional neural networks, recurrent neural networks, transformers, graph neural networks (\rightarrow ML2).
- ▶ **Beyond classification:** Neural networks for regression, mixture density networks, structured prediction (\rightarrow ML2).

对抗性鲁棒性 Adversarial robustness: Adversarial training, model certification.

- ▶ **Explainability:** Extracting insights from a neural network, Clever Hans effect (later this semester).

Summary

在实际应用中，神经网络面临两个主要挑战：

- 如何优化它们
- 如何正则化它们以便获得良好的泛化能力。

优化问题可以通过计算误差函数的海森矩阵（Hessian）和条件数（Condition number）来进行分析研究。

引入边际损失（Margin losses）和扰动（Perturbations）可以获得良好泛化能力的模型。

Previous lecture:

- ▶ Neural networks derive their **high representation power** from a large number of simple interconnected components (neurons), and the gradient of a neural network w.r.t. its parameters can be efficiently extracted (using **backpropagation**).

Today's lecture:

- ▶ In practice, there are two major challenges with neural networks (1) how to **optimize** them, and (2) how to **regularize** them so that they generalize well.
- ▶ The problem of optimization can be studied analytically by computing the **Hessian** of the error function and the **condition number**.
- ▶ Well-generalizing models can be obtained by introducing **margin losses** and **perturbations**.